

LOUNES BENALI

DEEP LEARNING ENGINEER | SPEECH & MULTIMODAL AI | LARGE-SCALE TRAINING & INFERENCE



B-Lounes

lounesbenali.2@gmail.com

+33 7 45 47 76 14



[lounes.benali](https://www.linkedin.com/in/lounes.benali)



[Portfolio](#)

PROFILE

Speech and multimodal AI engineer working across large-scale data curation, distributed GPU training, model adaptation, rigorous evaluation, and low-latency inference. At MBZUAI's Institute of Foundation Models, co-developing multilingual ASR for dialectal and code-switched speech on Slurm-managed AMD MI210 and NVIDIA H200 clusters.

PROFESSIONAL EXPERIENCE

Speech AI Research Intern

Feb 2026–Aug 2026 (expected)

Institute of Foundation Models (IFM), MBZUAI

Paris, France

Open foundation-model research lab at MBZUAI.

- Core contributor in a four-person speech team co-developing multilingual ASR for dialectal and code-switched speech. The first checkpoint ranked first in CER and third in WER on the public Open Universal ASR Leaderboard—within 2 WER points of leaders Cohere and Audar in Arabic—while supporting diacritized transcription, dialect identification, and code-switching across French, English, Spanish, and other languages.
- Engineered a resilient audio-ingestion system that collected nearly 4M hours, using deduplicated leased SQLite queues, partitioned VPN pools, pipelined download/Opus encoding, tmpfs-first I/O with disk fallback, and remotely verified dataset-shard uploads.
- Built a multi-model curation pipeline over 220M speech segments (roughly 1M hours), combining dialect ID, acoustic-quality labels, ASR hypotheses, forced alignment, and token confidence to produce a compact, high-quality 17k-hour ASR corpus.
- Built and operated large-scale Slurm workflows for distributed training and fault-tolerant inference across AMD/ROCm and NVIDIA/CUDA GPU clusters.
- Engineered a WebSocket streaming gateway around a non-streaming ASR model, combining causal VAD, language-aware hypothesis stabilization, final-first scheduling, and failure-aware endpoint leases; a 16-session replay achieved 1.33 s p95 first-partial and 0.55 s p95 post-EOS latency.
- Ran iterative self-supervised VietASR pretraining from random initialization plus supervised studies using as little as 100 labeled hours; scaled selected runs across 8–16 Slurm nodes and achieved 33.22 WER / 15.10 CER on a balanced 10-hour, 19-dialect evaluation set.
- Reimplemented a pinned FunASR FSMN-VAD frontend and causal encoder as batched PyTorch/ROCm operations, reaching 9,668× real-time end-to-end throughput on a single 8× AMD MI210 node after initialization.
- Ongoing work focuses on developing streaming-oriented audio encoders from scratch to support the broader IFM speech stack.

Deep Learning Engineer (Apprentice)

Sep 2024–Feb 2026

Ampere Software Technology (Renault Group)

Guyancourt, France

- Developed and fine-tuned TTS and speech-to-speech systems (Moshi), contributing from dataset generation and processing through model validation and deployment.
- Built an internal real-time inference platform serving multiple fine-tuned models, voices, versions, and languages; implemented streaming, model compression, inference acceleration, and latency instrumentation.
- Built a synthetic-conversation data pipeline that orchestrated two real-time voice agents, routed audio between them, converted sample rates, and recorded speaker-separated stereo WAVs for speech-model training.
- Built Kotlin/Android Automotive voice-assistant prototypes with coroutine-based WebSocket clients, full-duplex PCM audio, interruption control, jitter-aware playback, and pre-warmed session management.

Technical Staff – Paris 2024 Olympic Games

May 2024–Aug 2024

International Broadcast Center (IBC)

Le Bourget, France

- Deployed and supported computers, peripherals, and networked devices on the IBC LAN; diagnosed endpoint and connectivity issues during daily broadcast operations.

Electrical Engineering Intern

Jun 2023–Aug 2023

Schneider Electric

Algiers, Algeria

- Studied industrial variable-frequency-drive architectures and evaluated opportunities for machine-learning-based motor-control optimization.

SELECTED TECHNICAL WORK

- vLLM Audio-Model Integrations on ROCm – IFM, MBZUAI**

Python, PyTorch, vLLM, ROCm, Slurm

 - Implemented an out-of-tree HiggsAudio3Model plugin for vLLM, reproducing its chunked Whisper frontend and stride-2 projector, injecting audio embeddings into the native Qwen3 decoder, and exposing OpenAI-compatible transcription.
 - Replaced dynamically loaded remote configuration with importable classes compatible with vLLM EngineCore process serialization; validated greedy token/text parity against Transformers and deployed eight BF16 MI210 endpoints with Slurm lifecycle tooling.
 - Separately source-built vLLM for ROCm/gfx90a, integrated Audex/Nemotron-H, and brought up BF16 tensor-parallel audio serving across eight MI210 GPUs.
- Seenitt – Streaming Visual AI for Smart Glasses – MentraOS Grant**

Python, TensorRT, CUDA, [GitHub](#)

 - Received a MentraOS developer grant for Seenitt; built an Android-to-Python visual-AI stream that H.264-encodes camera frames on-device, decodes them with PyAV over WebSockets, and returns structured results.
 - Implemented a streaming chess-assistant prototype using SAM 3 to segment the physical board and pieces, reconstruct and track its 9 × 9 grid, stabilize observations against legal moves, emit FEN state, and return Stockfish suggestions.
 - Optimized the FP16 TensorRT path with cached text embeddings, reused vision features, page-locked buffers, and asynchronous CUDA transfers; overlapped GPU piece decoding with CPU grid reconstruction and tracked all 81 intersections using Lucas–Kanade optical flow.
 - Developing a vegetable-recognition prototype using prompt-driven segmentation, ViT classification, and best-view cataloguing, alongside a Blackjack companion prototype.
- k2 ROCm Port for Speech Training – IFM**

C++, HIP, PyTorch, ROCm, [GitHub](#)

 - Ported upstream k2’s CUDA-oriented C++/PyTorch extension to ROCm/HIP on AMD MI210, preserving the upstream API while adapting its build, packaging, and runtime dispatch.
 - Integrated PyTorch device/stream semantics, HIP/hipCUB compatibility, and kernel fixes for ragged-label prefix scans; validated representative FSA and autograd operations.
- Android Automotive Voice-AI Prototypes – Ampere**

Kotlin, Coroutines, WebSockets, [Hume](#) [GitHub](#)

 - Built Kotlin Android Automotive clients for hosted conversational-AI services, including Hume EVI, with 16-kHz PCM capture, buffered 24/48-kHz playback, live transcripts and prosody, interruption control, and tool-call flows.
 - Designed concurrency-safe pre-warmed session pools with atomic capacity control, readiness tracking, and language isolation; implemented RMS voice detection, jitter-buffer recovery, and phone, wired, Bluetooth, and automotive audio routing.

EDUCATION

- Master’s Degree in Intelligent Systems Engineering**

Sorbonne University

Sep 2024–Present

Paris, France
- Bachelor’s Degree in Electrical Engineering and Computer Science**

Paris-Saclay University

Sep 2023–May 2024

Paris, France

 - Graduated top of the class.
- Engineering Cycle – Electronics Engineering and Computer Science**

National Polytechnic School of Algiers

Sep 2022–Jun 2023

Algiers, Algeria
- Preparatory Classes in Science and Technology**

Two-year intensive national competitive-exam program

Sep 2020–Jun 2022

Algeria

 - Ranked 29th out of 1,576 candidates in the national competitive entrance examination for engineering schools.

TECHNICAL SKILLS

Speech & Multimodal AI: ASR, TTS, speech-to-speech, speech SSL, PyTorch, k2, Hugging Face

Distributed ML & Inference: Slurm, DDP, FSDP, vLLM, AMD/ROCm, NVIDIA/CUDA, multi-node training and inference

Vision & Edge Inference: SAM 3, TensorRT, ONNX Runtime, OpenCV, H.264/PyAV, optical flow

Programming: Python, C, C++, Kotlin, MATLAB, VHDL

Data & Systems: PostgreSQL, SQLite, DuckDB, Parquet, WebDataset, Docker, Git, WebSockets

Hardware & Embedded: FPGA/VHDL, Autodesk EAGLE PCB design

Technical Familiarity: Whisper-style encoders, Qwen AuT encoders, Conformer/FastConformer, streaming causal encoders, log-Mel/fbank frontends, continuous-audio versus tokenized-audio paradigms, RVQ-based audio tokens, Mimi codec, hybrid audio representations, audio-to-LLM adapters, Transformers, MoE, model parallelism, cross-modal attention, CTC, RNN-T, AED, LoRA, preference-based post-training (PPO/GRPO-style RLHF and DPO), diffusion, flow matching (FM, OT-CFM, radial-angular transport), FlashAttention, sliding-window attention, KV caching, PagedAttention, speculative decoding, continuous batching, pruning, quantization

LANGUAGES & ACTIVITIES

Languages: French (DALF C2); English and Arabic (fluent); Berber (native)

Interests: Robotics, Astronomy, Model Aircraft, Human–Machine Interfaces, Podcast Listening

Associative Life: Ex-IT Member at the IEEE Student Branch of the Polytechnic School of Algiers